

The following steps will show the prediction of diabetes in a person on the basis of different features like insulin, bmi,etc

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('diabetes.csv')
```

```
In [3]: df
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

observation:- from we can see that, the values of features are in different ranges. So we have to standardize them after train test split

```
In [4]: # check if standardization is required or not
df.std()
```

```
Out[4]:
```

	0
Pregnancies	3.369578
Glucose	31.972618
BloodPressure	19.355807
SkinThickness	15.952218
Insulin	115.244002
BMI	7.884160
DiabetesPedigreeFunction	0.331329
Age	11.760232
Outcome	0.476951

dtype: float64

from above it is clear that standardization is required

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Pregnancies         768 non-null   int64
1   Glucose             768 non-null   int64
2   BloodPressure       768 non-null   int64
3   SkinThickness       768 non-null   int64
4   Insulin             768 non-null   int64
5   BMI                 768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                 768 non-null   int64
8   Outcome             768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [6]: #check for null values
df.isnull().sum()
```

Out[6]:

	0
Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

dtype: int64  
no null values

```
In [7]: df['Age'].groupby(df['Outcome']).value_counts()
```

Out[7]:

		count
Outcome	Age	
0	22	61
	21	58
	24	38
	25	34
	23	31
...	...	...
1	67	1
	61	1
	55	1
	48	1
	70	1

96 rows × 1 columns

dtype: int64

```
In [8]: df.groupby('Outcome')['Age'].describe()
```

Out[8]:

	count	mean	std	min	25%	50%	75%	max
Outcome								
0	500.0	31.190000	11.667655	21.0	23.0	27.0	37.0	81.0
1	268.0	37.067164	10.968254	21.0	28.0	36.0	44.0	70.0

```
In [9]: df.groupby('Outcome').mean()
```

Out[9]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
Outcome								
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	0.429734	31.190000
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	0.550500	37.067164

this shows the mean of each feature regarding diabetic and none diabetic

```
In [10]: df.duplicated().sum()
```

Out[10]: 0

```
In [11]: df['Outcome'].value_counts()
```


```
Out[11]:
```

	count
Outcome	
0	500
1	268

dtype: int64

```
In [12]: sns.countplot(x='Outcome', data=df)
```

```
Out[12]: <Axes: xlabel='Outcome', ylabel='count'>
```



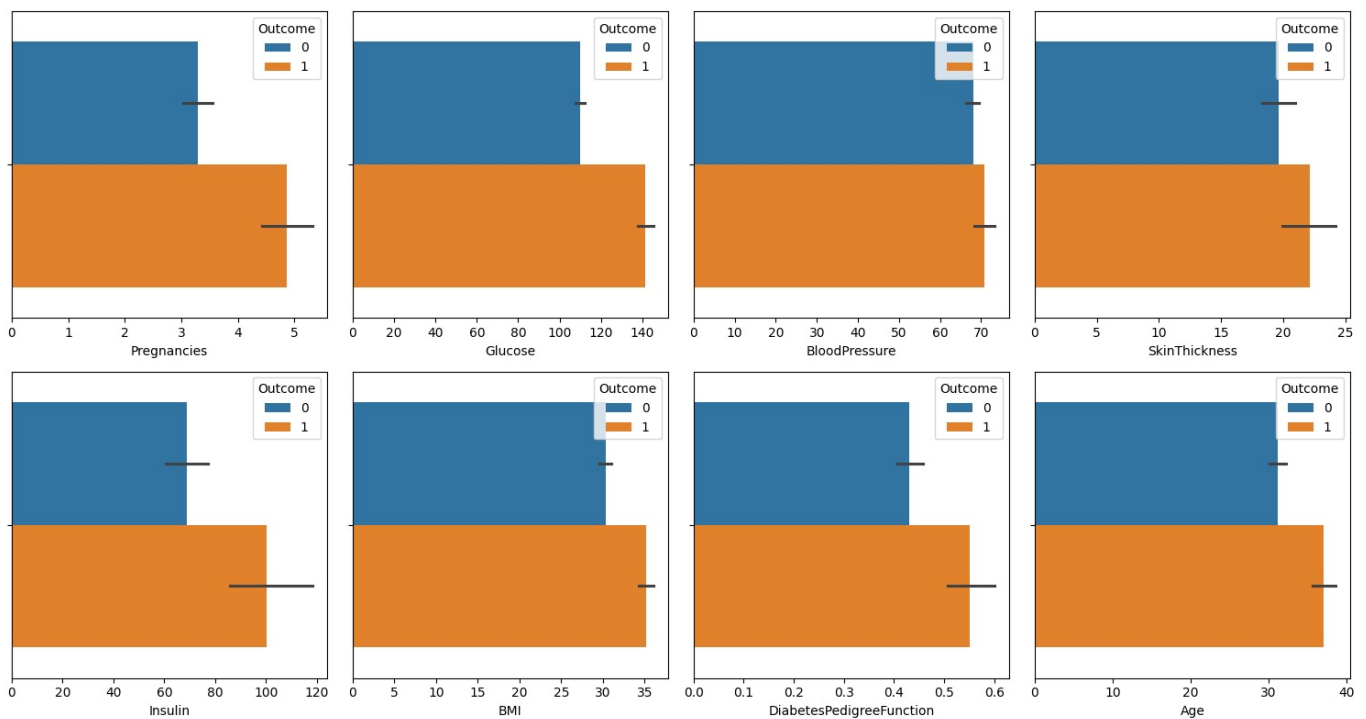
1--> Patient is Diabetic

0--> None Diabetic

to show the distribution of each feature with target feature(outcome)

```
In [13]: #create subplots
fig, axes = plt.subplots(nrows=2, ncols=4, figsize=(15, 8))

#plot distribution by iterating through the feature col
feature_col = df.columns.drop('Outcome')
for index, feature in enumerate(feature_col):
    row = index // 4 # Calculate the row index
    col = index % 4 # Calculate the column index
    sns.barplot(data=df, x=feature, hue='Outcome', ax=axes[row, col]) # Plot on the correct subplot
    plt.tight_layout()
plt.show()
```



now doing standardization

```
In [14]: from sklearn.preprocessing import StandardScaler

In [15]: scaler=StandardScaler()

In [16]: x=df.drop('Outcome',axis=1)

In [17]: y=df['Outcome']

In [18]: # train test split
from sklearn.model_selection import train_test_split

In [19]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,stratify=y,random_state=42)

In [20]: # scaling train and test data
x_train_scaled=scaler.fit_transform(x_train)
x_test_scaled=scaler.transform(x_test)
```

as the label is binary, we will be using classification models

And evaluating them to see which one works better in comparison to other for diabetics prediction

Used Classification models:-

1. Logistic Regression
2. SVM
3. Decision Tree
4. Random Forest
5. Naive Bayes

Logistic Regression

```
In [21]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,confusion_matrix,precision_score,recall_score,f1_score

In [22]: linear_model=LogisticRegression()

In [23]: linear_model.fit(x_train_scaled,y_train)

Out[23]: LogisticRegression
LogisticRegression()

In [24]: y_predict=linear_model.predict(x_test_scaled)
```

```
In [25]: log_accuracy=round(accuracy_score(y_test,y_predict),3)
log_accuracy
```

```
Out[25]: 0.714
```


```
In [27]: log_matrix=confusion_matrix(y_test,y_predict)
log_matrix
```

```
Out[27]: array([[82, 18],
               [26, 28]])
```

```
In [28]: log_precision=precision_score(y_test,y_predict)
log_recall=recall_score(y_test,y_predict)
log_f1=f1_score(y_test,y_predict)
```

```
In [29]: #visualize the confusion matrix
sns.heatmap(log_matrix,annot=True)
plt.ylabel('Actual')
plt.xlabel('Predicted')
```

```
Out[29]: Text(0.5, 23.52222222222222, 'Predicted')
```



Support Vector Machine

```
In [30]: from sklearn.svm import SVC
```

```
In [32]: machine=SVC(kernel='linear')
```

```
In [33]: machine.fit(x_train_scaled,y_train)
```

```
Out[33]: SVC(kernel='linear')
```

```
In [34]: y_predict=machine.predict(x_test_scaled)
machine_accuracy=round(accuracy_score(y_test,y_predict),3)
```

```
In [35]: machine_matrix=confusion_matrix(y_test,y_predict)
machine_precision=precision_score(y_test,y_predict)
machine_recall=recall_score(y_test,y_predict)
machine_f1=f1_score(y_test,y_predict)
print(machine_matrix)
```

```
[[83 17]
 [26 28]]
```

Decision Tree

```
In [36]: from sklearn.tree import DecisionTreeClassifier
```

```
In [37]: decision_model=DecisionTreeClassifier()
```

```
In [38]: decision_model.fit(x_train_scaled,y_train)
```

```
Out[38]: ▼ DecisionTreeClassifier
```

```
DecisionTreeClassifier()
```

```
In [39]: y_predict=decision_model.predict(x_test_scaled)
decision_accuracy=round(accuracy_score(y_test,y_predict),3)
```

```
In [40]: decision_matrix=confusion_matrix(y_test,y_predict)
decision_precision=precision_score(y_test,y_predict)
decision_recall=recall_score(y_test,y_predict)
decision_f1=f1_score(y_test,y_predict)
print(decision_matrix)
```

```
[[87 13]
 [31 23]]
```

Random Forest

```
In [41]: from sklearn.ensemble import RandomForestClassifier
```

```
In [42]: random_model=RandomForestClassifier()
```

```
In [43]: random_model.fit(x_train_scaled,y_train)
```

```
Out[43]: ▼ RandomForestClassifier
```

```
RandomForestClassifier()
```

```
In [44]: y_predict=random_model.predict(x_test_scaled)
random_accuracy=round(accuracy_score(y_test,y_predict),3)
```

```
In [45]: random_matrix=confusion_matrix(y_test,y_predict)
random_precision=precision_score(y_test,y_predict)
random_recall=recall_score(y_test,y_predict)
random_f1=f1_score(y_test,y_predict)
print(random_matrix)
```

```
[[81 19]
 [24 30]]
```

Naive Bayes

```
In [46]: from sklearn.naive_bayes import BernoulliNB
naive_model=BernoulliNB()
```

```
In [47]: naive_model.fit(x_train_scaled,y_train)
```

```
Out[47]: ▼ BernoulliNB
```

```
BernoulliNB()
```

```
In [48]: y_predict=naive_model.predict(x_test_scaled)
naive_accuracy=round(accuracy_score(y_test,y_predict),3)
```

```
In [49]: naive_matrix=confusion_matrix(y_test,y_predict)
naive_precision=precision_score(y_test,y_predict)
naive_recall=recall_score(y_test,y_predict)
naive_f1=f1_score(y_test,y_predict)
print(naive_matrix)
```

```
[[80 20]
 [26 28]]
```

```
In [50]: #comparision table
data={'model':['logistic reg','svm','decision tree','random forest','naive bayes'],
      'precision':[log_precision,machine_precision,decision_precision,random_precision,naive_precision],
      'recall':[log_recall,machine_recall,decision_recall,random_recall,naive_recall],
      'f1':[log_f1,machine_f1,decision_f1,random_f1,naive_f1],
      'accuracy':[log_accuracy,machine_accuracy,decision_accuracy,random_accuracy,naive_accuracy]}
comparison_df=pd.DataFrame(data)
```

```
In [51]: comparison_df
```

```
Out[51]:
```

	model	precision	recall	f1	accuracy
0	logistic reg	0.608696	0.518519	0.560000	0.714
1	svm	0.622222	0.518519	0.565657	0.721
2	decision tree	0.638889	0.425926	0.511111	0.714
3	random forest	0.612245	0.555556	0.582524	0.721
4	naive bayes	0.583333	0.518519	0.549020	0.701

as false negative(missing the diabetes) is more crucial --> recall should be more for less FN(priority)

as false positive(incorrectly diagnosing a disease) is less crucial than FN --> precision should be more for less FP

```
In [52]: machine_matrix
```

```
Out[52]: array([[83, 17],
               [26, 28]])
```

```
In [53]: random_matrix
```

```
Out[53]: array([[81, 19],
               [24, 30]])
```

Evaluation Metric -->

as F1-score balances Precision and Recall, we would be considering F1-Score for model evaluation

So according to that Random Forest would be best fit for Prediction of diabetes in a patient

```
In [54]: # define hyperparameter grid
from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

```
In [55]: # initialize gridsearch
grid_search = GridSearchCV(estimator=RandomForestClassifier(random_state=42),
                           param_grid=param_grid,
                           cv=5,
                           scoring='f1')
```

```
In [56]: #fit model
grid_search.fit(x_train_scaled, y_train)
```

```
Out[56]:
```

► **GridSearchCV**

► estimator: **RandomForestClassifier**

► RandomForestClassifier

```
In [58]: #best parameters and model
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
print(best_params)
print(best_model)

{'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 100}
RandomForestClassifier(min_samples_leaf=2, random_state=42)
```

```
In [59]: # evaluate the best model
y_pred = best_model.predict(x_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.7402597402597403

after tuning the accuracy of random forest model has increased from 0.72 to 0.74

```
In [60]: f1_value=f1_score(y_test,y_pred)
f1_value
```

```
Out[60]: 0.6
```

```
In [61]: # predictive system
```

```
input_data=(6,104,74,18,156,29.9,0.722,41)
```

```
#converting to array
```

```
data_array=np.asarray(input_data)
```

```
#reshapping
```

```
reshaped_data=data_array.reshape(1,-1)
```

```
#standardize the data as data would be in raw form and need to be standardized
```

```
data=scaler.transform(reshaped_data)
```

```
#predicting
```

```
prediction=random_model.predict(data)
```

```
#checking
```

```
if prediction[0]==0:
```

```
    print('patient is not diabetic')
```

```
else:
```

```
    print('patient is diabetic')
```

```
patient is diabetic
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
```

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js